

# Conways game of life in CUDA

## Introduction

This project implements **Conway's Game of Life** using **CUDA** to demonstrate **GPU** parallelism for fast and large-scale simulations on **Google collab**. The grid is initialized with a dense random pattern, filling 60% of the screen with live cells to create a dynamic evolution. By running on a **512×512 grid** over **500 iterations**.

This project essentially lets the user understand how to leverage **parallelism** for tasks that would otherwise consume a lot more time using only the CPU.

## Conways game of life

Conways game of life is a **cellular automaton**; in simple words it's a game with 4 rules. Those rules are:

1. Any live cell with **fewer than two live neighbors dies**, as if by underpopulation.
2. Any live cell with **two or three live neighbors lives** on to the next generation.
3. Any live cell with **more than three live neighbors dies**, as if by overpopulation.
4. Any dead cell with **exactly three live neighbors becomes a live cell**, as if by reproduction.

It's easier as well as more intuitive to understand its working by giving it a try. Click [here](#) to try it online.

The reason why chose Conways game of life for CUDA is that there are **multiple calculations** that need to be done, however these are relatively simple calculations and don't need a lot of power. This makes it ideal to run it on an GPU.

## CUDA

**CUDA (Compute Unified Device Architecture)** is the tool that lets us run our code on GPU's. Unlike traditional CPU programming, which executes instructions sequentially, CUDA enables thousands of **lightweight threads** to run in parallel. Click [here](#) to watch a 3-minute video explaining CUDA in detail.

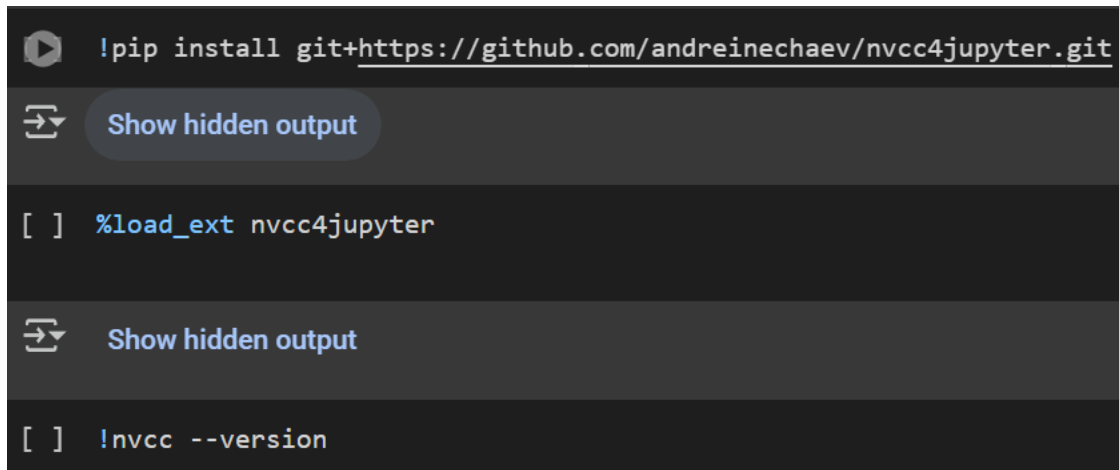
The key components of CUDA that we will be using in our project are:

1. Kernels
2. Grids
3. Blocks
4. Shared memory

## Code

We will be running the Entire code on Google Collab, making it accessible to everyone without requiring Actual GPU on your laptop!!

We install the prerequisite library



```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

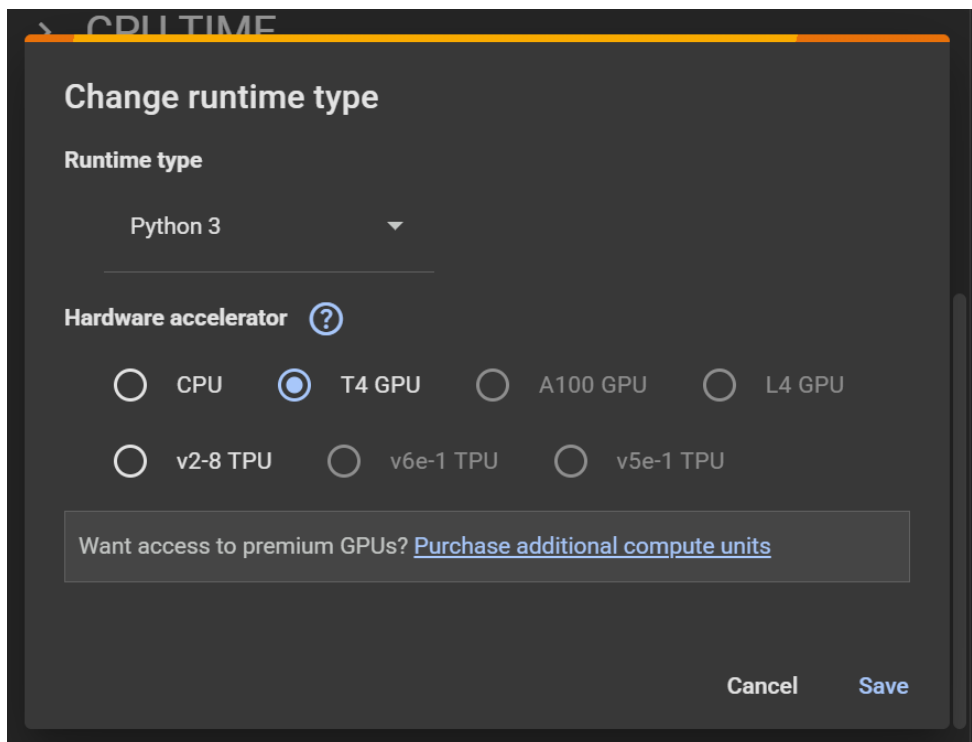
Show hidden output

```
[ ] %load_ext nvcc4jupyter
```

Show hidden output

```
[ ] !nvcc --version
```

Before, we run the code, we need to ensure that the we are connected to the T4 GPU. By default it connects to a CPU and won't work as intended.



# The GPU code

```
%%writefile buu.cpp
#include <iostream>
#include <cuda_runtime.h>
#include <chrono>
#include <random>

#define WIDTH 512
#define HEIGHT 512
#define ITERATIONS 1000

__device__ int count_neighbors(int* grid, int x, int y, int width, int height) {
    int count = 0;
    for (int dx = -1; dx <= 1; ++dx) {
        for (int dy = -1; dy <= 1; ++dy) {
            if (dx == 0 && dy == 0) continue;
            int nx = (x + dx + width) % width;
            int ny = (y + dy + height) % height;
            count += grid[ny * width + nx];
        }
    }
    return count;
}

__global__ void game_of_life_step(int* current, int* next, int width, int height) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;

    if (x < width && y < height) {
        int idx = y * width + x;
        int neighbors = count_neighbors(current, x, y, width, height);
        int state = current[idx];
        if (state == 1 && (neighbors == 2 || neighbors == 3)) {
            next[idx] = 1;
        } else if (state == 0 && neighbors == 3) {
            next[idx] = 1;
        } else {
            next[idx] = 0;
        }
    }
}

void set_random_dense_pattern(int* grid, int width, int height, float fill_ratio = 0.6f) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::bernoulli_distribution d(fill_ratio);
}
```

```

int main() {
    int size = WIDTH * HEIGHT * sizeof(int);

    int* h_grid = new int[WIDTH * HEIGHT];
    int* h_result = new int[WIDTH * HEIGHT];

    // Fill about 60% of the screen with live cells
    set_random_dense_pattern(h_grid, WIDTH, HEIGHT, 0.6f);

    int* d_current;
    int* d_next;
    cudaMalloc(&d_current, size);
    cudaMalloc(&d_next, size);
    cudaMemcpy(d_current, h_grid, size, cudaMemcpyHostToDevice);

    dim3 threadsPerBlock(16, 16);
    dim3 numBlocks((WIDTH + 15) / 16, (HEIGHT + 15) / 16);

    // CUDA timing setup
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);

    // GPU timing start
    auto gpu_start = std::chrono::high_resolution_clock::now();

    cudaEventRecord(start); // GPU timing start

    for (int i = 0; i < ITERATIONS; ++i) {
        game_of_life_step<<<numBlocks, threadsPerBlock>>>(d_current, d_next, WIDTH, HEIGHT);
        cudaDeviceSynchronize(); // Wait for kernel to finish

        std::swap(d_current, d_next);
        cudaMemcpy(h_result, d_current, size, cudaMemcpyDeviceToHost);
        //print_grid(h_result, WIDTH, HEIGHT);
    }

    cudaEventRecord(stop); // GPU timing stop
    cudaEventSynchronize(stop); // Ensure stop event completed

    // GPU timing end
    auto gpu_end = std::chrono::high_resolution_clock::now();
    std::chrono::duration<double, std::milli> gpu_duration = gpu_end - gpu_start;

    // Calculate GPU time
    // Output timing
    std::cout << "GPU kernel time: " << gpu_duration.count() << " ms\n";
}

```

# The CPU Code

```
[ ] %%writefile buu_cpu.cpp
#include <iostream>
#include <chrono>
#include <random>

#define WIDTH 512
#define HEIGHT 512
#define ITERATIONS 1000

void print_grid(int* grid) {
    for (int y = 0; y < HEIGHT; ++y) {
        for (int x = 0; x < WIDTH; ++x) {
            std::cout << (grid[y * WIDTH + x] ? '0' : '.');
        }
        std::cout << '\n';
    }
    std::cout << std::string(WIDTH, '=') << '\n';
}

void set_random_dense_pattern(int* grid, float fill_ratio = 0.6f) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::bernoulli_distribution d(fill_ratio);

    for (int y = 0; y < HEIGHT; ++y) {
        for (int x = 0; x < WIDTH; ++x) {
            grid[y * WIDTH + x] = d(gen) ? 1 : 0;
        }
    }
}
```

```

[ ] int main() {
    int size = WIDTH * HEIGHT * sizeof(int);

    int* current = new int[WIDTH * HEIGHT]();
    int* next = new int[WIDTH * HEIGHT]();

    // Fill about 60% of the screen with live cells
    set_random_dense_pattern(current, 0.6f);

    auto cpu_start = std::chrono::high_resolution_clock::now();

    for (int i = 0; i < ITERATIONS; ++i) {
        for (int y = 0; y < HEIGHT; y++) {
            for (int x = 0; x < WIDTH; x++) {

                int idx = y * WIDTH + x;

                int neighbors = 0;
                for (int dx = -1; dx <= 1; ++dx) {
                    for (int dy = -1; dy <= 1; ++dy) {
                        if (dx == 0 && dy == 0) continue;
                        int nx = (x + dx + WIDTH) % WIDTH;
                        int ny = (y + dy + HEIGHT) % HEIGHT;
                        neighbors += current[ny * WIDTH + nx];
                    }
                }

                int state = current[idx];
                if (state == 1 && (neighbors == 2 || neighbors == 3)) {
                    next[idx] = 1;
                } else if (state == 0 && neighbors == 3) {
                    next[idx] = 1;
                } else {
                    next[idx] = 0;
                }
            }
        }
    }
}

```

```

    int * temp = current;
    current = next;
    next = temp;

}

// CPU timing end
auto cpu_end = std::chrono::high_resolution_clock::now();
std::chrono::duration<double, std::milli> cpu_duration = cpu_end - cpu_start;


// Output timing
std::cout << "CPU total time: " << cpu_duration.count() << " ms" << std::endl;

// Cleanup
// cudaEventDestroy(start);
// cudaEventDestroy(stop);
// cudaFree(current);
// cudaFree(next);
delete[] current;
delete[] next;


return 0;
}

```

## Animation snippet:



## Results:



In the above test, we keep the grid size same at 512x512 and run the code for several iterations to compare the results.

We see that the GPU execution is **37** times faster than the CPU execution time (1000 iterations). In general for more iterations the GPU performs exponentially better when compared to the CPU.

## References

1. [Conways game of life-wikipedia](#)
2. [Numberphile-GOF](#)
3. [Computerphile-CUDA](#)
4. [Intro to CUDA](#)
5. [Nvidia-CUDA](#)